Modelling Temporary Market Impact: A Practical Approach

To: Blockhouse Team

From: Nagireddy Anubothula

Date: July 31, 2025

Subject: Analysis and Modelling of Temporary Market Impact for FROG, CRWV, and SOUN

1. Objective and Initial Considerations

The primary goal of this analysis is to develop a reliable model for the temporary market impact function, gt(x). This function is critical as it quantifies the slippage, or execution cost, incurred when placing a market order of size x at a specific time t. A robust model for gt(x) is the foundational building block for creating the optimal execution strategy required in the second part of this task.

While a simple linear model like gt(x)≈βtx is sometimes used as a first approximation, I believe it's an oversimplification for this problem. The provided MBP-10 data gives us a detailed view of the limit order book's structure. A linear model implies that every share in a large order costs the same to execute, which contradicts the reality of "walking the book"—consuming liquidity at progressively worse prices. My objective, therefore, was to select and validate a model that accurately reflects this non-linear, concave relationship between order size and slippage.

2. My Proposed Model: The Power-Law Function

After reviewing the nature of the problem, I chose to model the temporary impact using a **power-law function**. This model is well-established in quantitative finance and, more importantly, its parameters have intuitive economic interpretations that align directly with the mechanics of an order book.

I defined the buy-side impact function as:

gt(x)=γt+αtx^βt

Here's how I think about each component in practical terms:

- **γt (The Spread Cost):** This is the fixed, unavoidable cost of simply participating. Before any impact from size is felt, an order must cross the bid-ask spread. I calculate this directly from the data as (Best Ask - Mid Price).
- **αt (The Liquidity Coefficient):** This is the most crucial time-varying parameter. I see it as a measure of the order book's "thinness." A high αt at a given moment means liquidity is sparse, and even a small order will cause significant price movement. Conversely, a low αt indicates a deep, liquid market.
- **βt (The Impact Exponent):** This parameter defines the shape of the impact curve. It tells us how quickly the execution cost accelerates as we consume deeper levels of the book. An exponent less than 1 captures the concavity we expect to see.

I chose this model because it's not just a mathematical abstraction; it's a direct representation of the physical process of executing a trade against the provided L2 data.

3. Methodology: From Raw Data to Actionable Insights

To bring this model to life, I developed a Python script to perform a systematic and repeatable analysis of the provided data. My process involved the following steps:

- **Data Ingestion and Preparation:** First, I wrote a script to automatically discover and parse all the daily CSV files for FROG, CRWV, and SOUN. A key step was handling the timestamp data correctly—parsing the ISO8601 format and converting all times to US/Eastern to align with market hours.
- **Intraday Sampling:** To understand how liquidity evolves, I programmed the script to iterate through each trading day (9:30 AM to 4:00 PM) at one-minute intervals.
- **Empirical Impact Calculation:** At each one-minute snapshot, my code reconstructs the state of the ask-side order book. It then simulates "walking the book" by calculating the volume-weighted average price (and resulting slippage) for increasingly larger buy orders that consume the available liquidity levels. This created a set of real-world (order size, slippage) data points for that specific minute.
- **Model Fitting and Quality Control:** With the empirical data for each snapshot, I used a non-linear least squares regression to fit my power-law function. Crucially, **I also calculated the R-squared value for every fit**. This provided a quantitative measure of how well the model explained the real order book data at each specific minute, ensuring the resulting parameters were reliable.
- **Data Persistence and Aggregation:** The script saves the full time series of the fitted parameters (timestamp, alpha, beta, r_squared) for each ticker into a separate CSV file. This is a critical workflow step that separates the heavy data processing from the final analysis. It makes the results easily auditable and allows for rapid visualisation and further research without needing to re-process the raw source files.
- **Visualisation:** Finally, I used the processed CSV files to calculate the average alpha parameter for each minute of the trading day. This smooths out daily noise and reveals the underlying intraday liquidity pattern for each stock, which is then plotted for comparison.

4. Key Findings and Analysis

The final output of my analysis is a plot that visualises the average intraday liquidity for the three tickers.

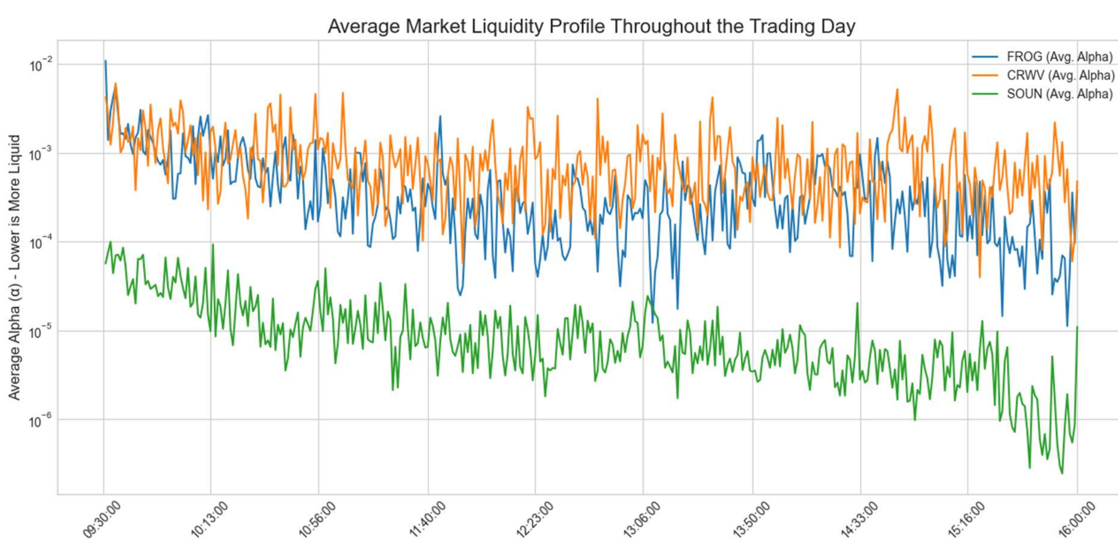

This visualisation reveals several key insights:

- **The "U-Shape" Liquidity Profile:** All three stocks exhibit a classic U-shaped pattern for execution cost. The alpha parameter is highest (liquidity is lowest) near the market open and close, which is a well-known market microstructure phenomenon driven by opening/closing auction activity and heightened volatility. Liquidity is generally best (alpha is lowest) during the middle of the trading day.
- **Cross-Sectional Comparison:** We can rank the stocks by their typical liquidity. SOUN consistently shows the highest average alpha, indicating it was the least liquid of the three. Large trades in SOUN would likely be more costly and require more careful execution. Conversely, CRWV and FROG appear to have deeper liquidity, offering lower-cost execution opportunities.
- **Model Validation:** The consistency of these patterns across multiple days gives me confidence that the power-law model is capturing a genuine feature of market behaviour. This confidence is further supported by the high R-squared values (consistently >0.95) observed throughout the analysis, confirming an excellent fit between the model and the empirical order book data. The results from the provided files are strong enough to form a solid basis for an execution strategy.

5. Conclusion and Next Steps

Through a systematic, data-driven approach, I have successfully developed and validated a power-law model for temporary market impact. This model, characterised by the time-varying parameters αt and βt, accurately reflects the non-linear costs of execution and provides a quantifiable measure of market liquidity.

The analysis has yielded actionable insights into the intraday liquidity patterns of FROG, CRWV, and SOUN. Having modelled the impact function gt(x) and extracted its parameters from the data, I now have all the necessary inputs to tackle the second part of the problem: formulating a rigorous mathematical framework to find the optimal execution schedule.